

# model\_training

March 17, 2019

## 1 Follow-Me Project

Congratulations on reaching the final project of the Robotics Nanodegree!

Previously, you worked on the Semantic Segmentation lab where you built a deep learning network that locates a particular human target within an image. For this project, you will utilize what you implemented and learned from that lab and extend it to train a deep learning model that will allow a simulated quadcopter to follow around the person that it detects!

Most of the code below is similar to the lab with some minor modifications. You can start with your existing solution, and modify and improve upon it to train the best possible model for this task.

You can click on any of the following to quickly jump to that part of this notebook: 1. Section ?? 2. Section ?? 3. Section ?? 4. Section [1.4](#) 5. Section [1.5](#) 6. Section [1.6](#)

### 1.1 Data Collection

We have provided you with a starting dataset for this project. Download instructions can be found in the README for this project's repo. Alternatively, you can collect additional data of your own to improve your model. Check out the "Collecting Data" section in the Project Lesson in the Classroom for more details!

```
In [1]: import os
import glob
import sys
import tensorflow as tf

from scipy import misc
import numpy as np

from tensorflow.contrib.keras.python import keras
from tensorflow.contrib.keras.python.keras import layers, models

from tensorflow import image

from utils import scoring_utils
from utils.separable_conv2d import SeparableConv2DKeras, BilinearUpSampling2D
from utils import data_iterator
from utils import plotting_tools
from utils import model_tools
```

## 1.2 FCN Layers

In the Classroom, we discussed the different layers that constitute a fully convolutional network (FCN). The following code will introduce you to the functions that you need to build your semantic segmentation model.

### 1.2.1 Separable Convolutions

The Encoder for your FCN will essentially require separable convolution layers, due to their advantages as explained in the classroom. The 1x1 convolution layer in the FCN, however, is a regular convolution. Implementations for both are provided below for your use. Each includes batch normalization with the ReLU activation function applied to the layers.

```
In [2]: def separable_conv2d_batchnorm(input_layer, filters, strides=1):
        output_layer = SeparableConv2DKeras(filters=filters, kernel_size=3, strides=strides,
                                             padding='same', activation='relu')(input_layer)

        output_layer = layers.BatchNormalization()(output_layer)
        return output_layer

def conv2d_batchnorm(input_layer, filters, kernel_size=3, strides=1):
    output_layer = layers.Conv2D(filters=filters, kernel_size=kernel_size, strides=strides,
                                 padding='same', activation='relu')(input_layer)

    output_layer = layers.BatchNormalization()(output_layer)
    return output_layer
```

### 1.2.2 Bilinear Upsampling

The following helper function implements the bilinear upsampling layer. Upsampling by a factor of 2 is generally recommended, but you can try out different factors as well. Upsampling is used in the decoder block of the FCN.

```
In [3]: def bilinear_upsample(input_layer):
        output_layer = BilinearUpSampling2D((2,2))(input_layer)
        return output_layer
```

## 1.3 Build the Model

In the following cells, you will build an FCN to train a model to detect and locate the hero target within an image. The steps are: - Create an `encoder_block` - Create a `decoder_block` - Build the FCN consisting of encoder block(s), a 1x1 convolution, and decoder block(s). This step requires experimentation with different numbers of layers and filter sizes to build your model.

### 1.3.1 Encoder Block

Create an encoder block that includes a separable convolution layer using the `separable_conv2d_batchnorm()` function. The `filters` parameter defines the size or depth of the output layer. For example, 32 or 64.

```
In [4]: def encoder_block(input_layer, filters, strides):
```

```
    # TODO Create a separable convolution layer using the separable_conv2d_batchnorm() f
    output_layer = separable_conv2d_batchnorm(input_layer, filters, strides)
    return output_layer
```

### 1.3.2 Decoder Block

The decoder block is comprised of three parts: - A bilinear upsampling layer using the `upsample_bilinear()` function. The current recommended factor for upsampling is set to 2. - A layer concatenation step. This step is similar to skip connections. You will concatenate the upsampled `small_ip_layer` and the `large_ip_layer`. - Some (one or two) additional separable convolution layers to extract some more spatial information from prior layers.

```
In [5]: def decoder_block(small_ip_layer, large_ip_layer, filters):
```

```
    # TODO Upsample the small input layer using the bilinear_upsample() function.
    upsampled_layer = bilinear_upsample(small_ip_layer)
    print("upsampled-shape:", upsampled_layer.get_shape())
    # TODO Concatenate the upsampled and large input layers using layers.concatenate
    print("large_ip_layer-shape", large_ip_layer.get_shape())
    concatenate_layers = layers.concatenate([upsampled_layer, large_ip_layer])
    print("concatenate-shape:", concatenate_layers.get_shape())
    # TODO Add some number of separable convolution layers
    output_layer = separable_conv2d_batchnorm(concatenate_layers, filters, strides = 1)
    print("decoder_output-shape:", output_layer.get_shape())
    return output_layer

    return output_layer
```

### 1.3.3 Model

Now that you have the encoder and decoder blocks ready, go ahead and build your FCN architecture!

There are three steps: - Add encoder blocks to build the encoder layers. This is similar to how you added regular convolutional layers in your CNN lab. - Add a 1x1 Convolution layer using the `conv2d_batchnorm()` function. Remember that 1x1 Convolutions require a kernel and stride of 1. - Add decoder blocks for the decoder layers.

```
In [6]: def fcn_model(inputs, num_classes):
```

```
    # TODO Add Encoder Blocks.
    # Remember that with each encoder layer, the depth of your model (the number of filt
    encoder_1 = encoder_block(inputs, 64, strides = 2)
    print("encoder_1-shape:", encoder_1.get_shape())
    encoder_2 = encoder_block(encoder_1, 128, strides = 2)
    print("encoder_2-shape:", encoder_2.get_shape())
    # TODO Add 1x1 Convolution layer using conv2d_batchnorm().
    conv_1x1 = conv2d_batchnorm(encoder_2, filters = 256, kernel_size=1, strides=1)
```

```

print("conv_1x1-shape:", conv_1x1.get_shape())
# TODO: Add the same number of Decoder Blocks as the number of Encoder Blocks
decoder_1 = decoder_block(conv_1x1, encoder_1, 64)
decoder_2 = decoder_block(decoder_1, inputs, 32)
# The function returns the output layer of your model. "x" is the final layer obtained
return layers.Conv2D(num_classes, 3, activation='softmax', padding='same')(decoder_2)

```

## 1.4 Training

The following cells will use the FCN you created and define an output layer based on the size of the processed image and the number of classes recognized. You will define the hyperparameters to compile and train your model.

Please Note: For this project, the helper code in `data_iterator.py` will resize the copter images to 160x160x3 to speed up training.

```

In [7]: """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        image_hw = 160
        image_shape = (image_hw, image_hw, 3)
        inputs = layers.Input(image_shape)
        num_classes = 3

        # Call fcn_model()
        output_layer = fcn_model(inputs, num_classes)
        print("output-shape:", output_layer.get_shape())

```

```

encoder_1-shape: (?, 80, 80, 64)
encoder_2-shape: (?, 40, 40, 128)
conv_1x1-shape: (?, 40, 40, 256)
upsampled-shape: (?, 80, 80, 256)
large_ip_layer-shape (?, 80, 80, 64)
concatenate-shape: (?, 80, 80, 320)
decoder_output-shape: (?, 80, 80, 64)
upsampled-shape: (?, 160, 160, 64)
large_ip_layer-shape (?, 160, 160, 3)
concatenate-shape: (?, 160, 160, 67)
decoder_output-shape: (?, 160, 160, 32)
output-shape: (?, 160, 160, 3)

```

### 1.4.1 Hyperparameters

Define and tune your hyperparameters. - **batch\_size**: number of training samples/images that get propagated through the network in a single pass. - **num\_epochs**: number of times the entire training dataset gets propagated through the network. - **steps\_per\_epoch**: number of batches of

training images that go through the network in 1 epoch. We have provided you with a default value. One recommended value to try would be based on the total number of images in training dataset divided by the batch\_size. - **validation\_steps**: number of batches of validation images that go through the network in 1 epoch. This is similar to steps\_per\_epoch, except validation\_steps is for the validation dataset. We have provided you with a default value for this as well. - **workers**: maximum number of processes to spin up. This can affect your training speed and is dependent on your hardware. We have provided a recommended value to work with.

```
In [8]: learning_rate = 0.001
        batch_size = 32
        num_epochs = 100
        steps_per_epoch = 200
        validation_steps = 50
        workers = 4
```

```
In [9]: """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        from workspace_utils import active_session
        # Keeping Your Session Active
        with active_session():
            # Define the Keras model and compile it for training
            model = models.Model(inputs=inputs, outputs=output_layer)

            model.compile(optimizer=keras.optimizers.Adam(learning_rate), loss='categorical_crossentropy')

            # Data iterators for loading the training and validation data
            train_iter = data_iterator.BatchIteratorSimple(batch_size=batch_size,
                                                            data_folder=os.path.join('..', 'data'),
                                                            image_shape=image_shape,
                                                            shift_aug=True)


            val_iter = data_iterator.BatchIteratorSimple(batch_size=batch_size,
                                                         data_folder=os.path.join('..', 'data'),
                                                         image_shape=image_shape)

            logger_cb = plotting_tools.LoggerPlotter()
            callbacks = [logger_cb]

            model.fit_generator(train_iter,
                                steps_per_epoch = steps_per_epoch, # the number of batches per epoch
                                epochs = num_epochs, # the number of epochs to train for,
                                validation_data = val_iter, # validation iterator
                                validation_steps = validation_steps, # the number of batches to
                                callbacks=callbacks,
                                workers = workers)
```

Epoch 1/100


199/200 [=====>.] - ETA: 0s - loss: 0.3396



200/200 [=====] - 101s - loss: 0.3383 - val\_loss: 0.1114

Epoch 2/100


199/200 [=====>.] - ETA: 0s - loss: 0.0512




200/200 [=====] - 96s - loss: 0.0513 - val\_loss: 0.0987

Epoch 3/100

199/200 [=====>.] - ETA: 0s - loss: 0.0432




200/200 [=====] - 97s - loss: 0.0432 - val\_loss: 0.0547  
Epoch 4/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0393



200/200 [=====] - 96s - loss: 0.0393 - val\_loss: 0.0497  
Epoch 5/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0385







200/200 [=====] - 97s - loss: 0.0385 - val\_loss: 0.0493

Epoch 6/100


199/200 [=====>.] - ETA: 0s - loss: 0.0350



200/200 [=====] - 96s - loss: 0.0350 - val\_loss: 0.0357  
Epoch 7/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0332




200/200 [=====] - 97s - loss: 0.0332 - val\_loss: 0.0387  
Epoch 8/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0349




200/200 [=====] - 96s - loss: 0.0349 - val\_loss: 0.0645

Epoch 9/100


199/200 [=====>.] - ETA: 0s - loss: 0.0322



200/200 [=====] - 96s - loss: 0.0321 - val\_loss: 0.0345  
Epoch 10/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0301




200/200 [=====] - 97s - loss: 0.0301 - val\_loss: 0.0381  
Epoch 11/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0290




200/200 [=====] - 97s - loss: 0.0290 - val\_loss: 0.0372

Epoch 12/100


199/200 [=====>.] - ETA: 0s - loss: 0.0300



200/200 [=====] - 96s - loss: 0.0300 - val\_loss: 0.0316  
Epoch 13/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0274




200/200 [=====] - 96s - loss: 0.0274 - val\_loss: 0.0340  
Epoch 14/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0265




200/200 [=====] - 96s - loss: 0.0267 - val\_loss: 0.0383

Epoch 15/100

199/200 [=====>.] - ETA: 0s - loss: 0.0280




200/200 [=====] - 97s - loss: 0.0280 - val\_loss: 0.0381  
Epoch 16/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0262




200/200 [=====] - 97s - loss: 0.0262 - val\_loss: 0.0327  
Epoch 17/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0260







200/200 [=====] - 96s - loss: 0.0260 - val\_loss: 0.0347  
 Epoch 18/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0263



200/200 [=====] - 97s - loss: 0.0263 - val\_loss: 0.0363  
Epoch 19/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0239




200/200 [=====] - 96s - loss: 0.0239 - val\_loss: 0.0375  
Epoch 20/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0256




200/200 [=====] - 97s - loss: 0.0256 - val\_loss: 0.0311

Epoch 21/100


199/200 [=====>.] - ETA: 0s - loss: 0.0245



200/200 [=====] - 96s - loss: 0.0245 - val\_loss: 0.0356  
Epoch 22/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0246




200/200 [=====] - 96s - loss: 0.0248 - val\_loss: 0.0355  
Epoch 23/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0239




200/200 [=====] - 96s - loss: 0.0239 - val\_loss: 0.0326

Epoch 24/100


199/200 [=====>.] - ETA: 0s - loss: 0.0233



200/200 [=====] - 96s - loss: 0.0232 - val\_loss: 0.0304  
Epoch 25/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0230




200/200 [=====] - 97s - loss: 0.0230 - val\_loss: 0.0357  
Epoch 26/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0223




200/200 [=====] - 96s - loss: 0.0223 - val\_loss: 0.0258

Epoch 27/100

199/200 [=====>.] - ETA: 0s - loss: 0.0230




200/200 [=====] - 97s - loss: 0.0230 - val\_loss: 0.0315  
Epoch 28/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0214



200/200 [=====] - 97s - loss: 0.0213 - val\_loss: 0.0301  
Epoch 29/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0221







200/200 [=====] - 97s - loss: 0.0221 - val\_loss: 0.0295

Epoch 30/100


199/200 [=====>.] - ETA: 0s - loss: 0.0228



200/200 [=====] - 96s - loss: 0.0228 - val\_loss: 0.0370  
Epoch 31/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0234




200/200 [=====] - 97s - loss: 0.0233 - val\_loss: 0.0281  
Epoch 32/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0215




200/200 [=====] - 96s - loss: 0.0215 - val\_loss: 0.0324

Epoch 33/100


199/200 [=====>.] - ETA: 0s - loss: 0.0203



200/200 [=====] - 97s - loss: 0.0206 - val\_loss: 0.0414  
Epoch 34/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0214




200/200 [=====] - 96s - loss: 0.0215 - val\_loss: 0.0284  
Epoch 35/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0208




200/200 [=====] - 97s - loss: 0.0208 - val\_loss: 0.0306

Epoch 36/100


199/200 [=====>.] - ETA: 0s - loss: 0.0207



200/200 [=====] - 96s - loss: 0.0207 - val\_loss: 0.0280  
Epoch 37/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0204




200/200 [=====] - 97s - loss: 0.0204 - val\_loss: 0.0295  
Epoch 38/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0201




200/200 [=====] - 96s - loss: 0.0200 - val\_loss: 0.0272

Epoch 39/100

199/200 [=====>.] - ETA: 0s - loss: 0.0193




200/200 [=====] - 96s - loss: 0.0193 - val\_loss: 0.0290  
Epoch 40/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0264



200/200 [=====] - 96s - loss: 0.0263 - val\_loss: 0.0469  
Epoch 41/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0233







200/200 [=====] - 96s - loss: 0.0232 - val\_loss: 0.0289

Epoch 42/100


199/200 [=====>.] - ETA: 0s - loss: 0.0208



200/200 [=====] - 97s - loss: 0.0208 - val\_loss: 0.0277  
Epoch 43/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0189




200/200 [=====] - 96s - loss: 0.0189 - val\_loss: 0.0283  
Epoch 44/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0204




200/200 [=====] - 96s - loss: 0.0204 - val\_loss: 0.0317

Epoch 45/100


199/200 [=====>.] - ETA: 0s - loss: 0.0191



200/200 [=====] - 96s - loss: 0.0194 - val\_loss: 0.0290  
Epoch 46/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0183




200/200 [=====] - 97s - loss: 0.0183 - val\_loss: 0.0266  
Epoch 47/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0195




200/200 [=====] - 96s - loss: 0.0195 - val\_loss: 0.0336

Epoch 48/100


199/200 [=====>.] - ETA: 0s - loss: 0.0175




200/200 [=====] - 97s - loss: 0.0174 - val\_loss: 0.0267  
Epoch 49/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0200




200/200 [=====] - 96s - loss: 0.0200 - val\_loss: 0.0269  
Epoch 50/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0190



200/200 [=====] - 96s - loss: 0.0190 - val\_loss: 0.0256  
 Epoch 51/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0183




200/200 [=====] - 96s - loss: 0.0183 - val\_loss: 0.0255  
Epoch 52/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0174



200/200 [=====] - 96s - loss: 0.0174 - val\_loss: 0.0359  
Epoch 53/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0184







200/200 [=====] - 97s - loss: 0.0184 - val\_loss: 0.0323

Epoch 54/100


199/200 [=====>.] - ETA: 0s - loss: 0.0183



200/200 [=====] - 96s - loss: 0.0183 - val\_loss: 0.0276  
Epoch 55/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0176




200/200 [=====] - 97s - loss: 0.0176 - val\_loss: 0.0273  
Epoch 56/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0177




200/200 [=====] - 96s - loss: 0.0177 - val\_loss: 0.0269

Epoch 57/100


199/200 [=====>.] - ETA: 0s - loss: 0.0171



200/200 [=====] - 97s - loss: 0.0171 - val\_loss: 0.0294  
Epoch 58/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0178




200/200 [=====] - 97s - loss: 0.0178 - val\_loss: 0.0270  
Epoch 59/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0179




```

200/200 [=====] - 97s - loss: 0.0179 - val_loss: 0.0267
Epoch 60/100
199/200 [=====>.] - ETA: 0s - loss: 0.0181


```




200/200 [=====] - 96s - loss: 0.0181 - val\_loss: 0.0257  
Epoch 61/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0182




200/200 [=====] - 97s - loss: 0.0182 - val\_loss: 0.0304  
Epoch 62/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0201



200/200 [=====] - 97s - loss: 0.0201 - val\_loss: 0.0372  
 Epoch 63/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0189




200/200 [=====] - 97s - loss: 0.0189 - val\_loss: 0.0302  
Epoch 64/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0170




200/200 [=====] - 96s - loss: 0.0170 - val\_loss: 0.0270  
Epoch 65/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0177







200/200 [=====] - 96s - loss: 0.0177 - val\_loss: 0.0270  
 Epoch 66/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0162



200/200 [=====] - 97s - loss: 0.0163 - val\_loss: 0.0260  
Epoch 67/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0170




200/200 [=====] - 97s - loss: 0.0170 - val\_loss: 0.0272  
Epoch 68/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0160




200/200 [=====] - 96s - loss: 0.0160 - val\_loss: 0.0330

Epoch 69/100


199/200 [=====>.] - ETA: 0s - loss: 0.0188



200/200 [=====] - 96s - loss: 0.0188 - val\_loss: 0.0332  
Epoch 70/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0169




200/200 [=====] - 97s - loss: 0.0169 - val\_loss: 0.0297  
Epoch 71/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0157




```


200/200 [=====] - 96s - loss: 0.0157 - val_loss: 0.0234
Epoch 72/100
199/200 [=====>.] - ETA: 0s - loss: 0.0166
  
```




200/200 [=====] - 97s - loss: 0.0166 - val\_loss: 0.0285  
Epoch 73/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0161




200/200 [=====] - 97s - loss: 0.0161 - val\_loss: 0.0311  
Epoch 74/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0157



200/200 [=====] - 96s - loss: 0.0157 - val\_loss: 0.0258  
 Epoch 75/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0164




200/200 [=====] - 97s - loss: 0.0164 - val\_loss: 0.0315  
Epoch 76/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0163




200/200 [=====] - 96s - loss: 0.0163 - val\_loss: 0.0278  
Epoch 77/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0165







200/200 [=====] - 97s - loss: 0.0164 - val\_loss: 0.0314  
 Epoch 78/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0153




200/200 [=====] - 96s - loss: 0.0152 - val\_loss: 0.0268  
Epoch 79/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0164




200/200 [=====] - 96s - loss: 0.0165 - val\_loss: 0.0266  
Epoch 80/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0154




200/200 [=====] - 96s - loss: 0.0154 - val\_loss: 0.0320  
 Epoch 81/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0159




200/200 [=====] - 97s - loss: 0.0159 - val\_loss: 0.0337  
Epoch 82/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0154




200/200 [=====] - 96s - loss: 0.0154 - val\_loss: 0.0287  
Epoch 83/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0147




200/200 [=====] - 96s - loss: 0.0147 - val\_loss: 0.0290  
 Epoch 84/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0156



200/200 [=====] - 96s - loss: 0.0156 - val\_loss: 0.0274  
Epoch 85/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0152




200/200 [=====] - 96s - loss: 0.0152 - val\_loss: 0.0345  
Epoch 86/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0152




200/200 [=====] - 96s - loss: 0.0152 - val\_loss: 0.0276

Epoch 87/100

199/200 [=====>.] - ETA: 0s - loss: 0.0159




200/200 [=====] - 97s - loss: 0.0159 - val\_loss: 0.0281  
Epoch 88/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0153




200/200 [=====] - 96s - loss: 0.0153 - val\_loss: 0.0280  
Epoch 89/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0160







200/200 [=====] - 97s - loss: 0.0160 - val\_loss: 0.0274  
 Epoch 90/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0154




200/200 [=====] - 97s - loss: 0.0154 - val\_loss: 0.0288  
Epoch 91/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0150




200/200 [=====] - 96s - loss: 0.0150 - val\_loss: 0.0266  
Epoch 92/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0157




200/200 [=====] - 97s - loss: 0.0156 - val\_loss: 0.0307  
 Epoch 93/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0148




200/200 [=====] - 96s - loss: 0.0148 - val\_loss: 0.0534  
Epoch 94/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0181




200/200 [=====] - 96s - loss: 0.0180 - val\_loss: 0.0257  
Epoch 95/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0151




200/200 [=====] - 97s - loss: 0.0151 - val\_loss: 0.0268  
 Epoch 96/100  
 199/200 [=====>.] - ETA: 0s - loss: 0.0146



200/200 [=====] - 97s - loss: 0.0146 - val\_loss: 0.0274  
Epoch 97/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0181




200/200 [=====] - 96s - loss: 0.0180 - val\_loss: 0.0288  
Epoch 98/100  
199/200 [=====>.] - ETA: 0s - loss: 0.0145



200/200 [=====] - 96s - loss: 0.0145 - val\_loss: 0.0288

Epoch 99/100


199/200 [=====>.] - ETA: 0s - loss: 0.0141



```

200/200 [=====] - 96s - loss: 0.0141 - val_loss: 0.0289
Epoch 100/100
199/200 [=====>.] - ETA: 0s - loss: 0.0143

```



```

200/200 [=====] - 97s - loss: 0.0144 - val_loss: 0.0293

```

```

In [10]: # Save your trained model weights
         weight_file_name = 'model_weights'
         model_tools.save_network(model, weight_file_name)

```

## 1.5 Prediction

Now that you have your model trained and saved, you can make predictions on your validation dataset. These predictions can be compared to the mask images, which are the ground truth labels, to evaluate how well your model is doing under different conditions.

There are three different predictions available from the helper code provided: - **patrol\_with\_targ**: Test how well the network can detect the hero from a distance. - **patrol\_non\_targ**: Test how often the network makes a mistake and identifies the wrong person as the target. - **following\_images**: Test how well the network can identify the target while following them.



```
In [11]: # If you need to load a model which you previously trained you can uncomment the code below

# weight_file_name = 'model_weights'
# restored_model = model_tools.load_network(weight_file_name)
```

The following cell will write predictions to files and return paths to the appropriate directories. The run\_num parameter is used to define or group all the data for a particular model run. You can change it for different runs. For example, 'run\_1', 'run\_2' etc.

```
In [12]: run_num = 'run_1'


val_with_targ, pred_with_targ = model_tools.write_predictions_grade_set(model,
                                                                    run_num, 'patrol_with_targ', 'sample_evaluation_data')



val_no_targ, pred_no_targ = model_tools.write_predictions_grade_set(model,
                                                                    run_num, 'patrol_non_targ', 'sample_evaluation_data')

val_following, pred_following = model_tools.write_predictions_grade_set(model,
                                                                    run_num, 'following_images', 'sample_evaluation_data')
```


Now lets look at your predictions, and compare them to the ground truth labels and original images. Run each of the following cells to visualize some sample images from the predictions in the validation set.



```
In [17]: # images while following the target
im_files = plotting_tools.get_im_file_sample('sample_evaluation_data', 'following_images')
for i in range(3):
    im_tuple = plotting_tools.load_images(im_files[i])
    plotting_tools.show_images(im_tuple)
```






```
In [18]: # images while at patrol without target
im_files = plotting_tools.get_im_file_sample('sample_evaluation_data', 'patrol_non_targ'
for i in range(3):
    im_tuple = plotting_tools.load_images(im_files[i])
    plotting_tools.show_images(im_tuple)
```






In [19]:

```
# images while at patrol with target
im_files = plotting_tools.get_im_file_sample('sample_evaluation_data','patrol_with_targ
for i in range(3):
    im_tuple = plotting_tools.load_images(im_files[i])
    plotting_tools.show_images(im_tuple)
```





## 1.6 Evaluation

Evaluate your model! The following cells include several different scores to help you evaluate your model under the different conditions discussed during the Prediction step.

In [20]: *# Scores for while the quad is following behind the target.*

```
true_pos1, false_pos1, false_neg1, iou1 = scoring_utils.score_run_iou(val_following, pr
```

number of validation samples intersection over the union evaluated on 542

average intersection over union for background is 0.9952069229249436

average intersection over union for other people is 0.348246367183434

average intersection over union for the hero is 0.8749019619801721

number true positives: 539, number false positives: 0, number false negatives: 0

In [21]: *# Scores for images while the quad is on patrol and the target is not visible*

```
true_pos2, false_pos2, false_neg2, iou2 = scoring_utils.score_run_iou(val_no_targ, pred
```

```
number of validation samples intersection over the union evaulated on 270
average intersection over union for background is 0.9861884958680348
average intersection over union for other people is 0.721654260088537
average intersection over union for the hero is 0.0
number true positives: 0, number false positives: 62, number false negatives: 0
```

```
In [22]: # This score measures how well the neural network can detect the target from far away
        true_pos3, false_pos3, false_neg3, iou3 = scoring_utils.score_run_iou(val_with_targ, pr
```

```
number of validation samples intersection over the union evaulated on 322
average intersection over union for background is 0.9960621111155227
average intersection over union for other people is 0.43813107303103743
average intersection over union for the hero is 0.22699085853805873
number true positives: 146, number false positives: 1, number false negatives: 155
```

```
In [24]: # Sum all the true positives, etc from the three datasets to get a weight for the score
        true_pos = true_pos1 + true_pos2 + true_pos3
        false_pos = false_pos1 + false_pos2 + false_pos3
        false_neg = false_neg1 + false_neg2 + false_neg3

        weight = true_pos/(true_pos+false_neg+false_pos)
        print(weight)
```

```
0.7585825027685493
```

```
In [25]: # The IoU for the dataset that never includes the hero is excluded from grading
        final_IoU = (iou1 + iou3)/2
        print(final_IoU)
```

```
0.550946410259
```

```
In [26]: # And the final grade score is
        final_score = final_IoU * weight
        print(final_score)
```

```
0.417938306786
```

```
In [ ]:
```